

# AP GCC - CGA tool chain on unix systems

Torsten Becker, Frederik Rudeck, Robert Timm

2009-08-01

## Zusammenfassung

The CGA - Call Graph Analyzer - system developed at the HPI was indented for Windows based systems only. We ported this tool chain to Mac OS X and Linux systems.

## TEMPLATE

Inhaltsvorlage zur Dokumentation - Seminar Softwarevisualisierung SoSe09

Arbeitspaket: Linux/MacOS-Portierung von CGA

Portierung von CGA nach Linux/MacOS

Hindernisse bei der Portierung

todo

Coding Richtlinien zur zukünftigen Vermeidung von Cross-Plattform Problemen

Todo

Cross-Plattform Buildsystem CMake

Nutzung von CMake unter Windows, Linux und MacOS

todo

CGA Coding Richtlinien zur Nutzung von CMake

Todo

Portierung des Windows-basierten Faktextraktionsmechanismus Callmon nach Linux/MacOS

Analysierbare Softwaresysteme - Notwendige Voraussetzungen

Mit welcher GCC Version muss das zu analysierende System gebaut sein?

Gibt es bestimmte Systemlibs, gegen die das zu analysierende System gelinkt werden muss?

Sind Inkompatibilitäten zu erwarten mit irgendwelchen Systemlibs?

...

Todo

Workflow der Callmon Toolchain

Todo

Einbindung von 3rd Party Tools

Welche Datenquelle für Debug-Informationen wird verwendet? Gibt es Alternativen? Warum die eine gewählt?

Wie werden die Funktionseinsprungsadressen und Funktionsausstiegsadressen bestimmt? Gibt es Alternativen?

Implementierung der Callmon Bibliothek unter Linux/MacOS

Wie ist der Callmon-Kern nach Linux/MacOS portiert worden? Wie wurde der Start/Stop-Mechanismus übertragen?

Todo

Tutorial: Kontrollfluss-Analyse von Audacity

Charakterisierung des Softwaresystems (LOC, etc.)

Todo

Exemplarisches Vorgehen bei der Faktextraktion (und anschließenden Visualisierung)

Todo

# Inhaltsverzeichnis

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Challenges</b>	<b>4</b>
2.1	Complexity . . . . .	4
2.2	Platform specifics . . . . .	4
2.3	Compiler specifics . . . . .	4
2.4	IDE specifics . . . . .	4
<b>3</b>	<b>Results</b>	<b>5</b>
3.1	CMake build enviroment . . . . .	5
3.1.1	Qt specific build steps . . . . .	5
3.2	Callmon runtime library . . . . .	5
3.2.1	Filesystem events . . . . .	5
3.2.2	Lockfree event to disk . . . . .	5
3.2.3	Asynchronous IO . . . . .	5
3.3	Patch and Patchclean . . . . .	5
3.3.1	ELF patching . . . . .	5
3.3.2	MachO patching . . . . .	5
3.4	Metacreator . . . . .	5
3.4.1	Line number caching . . . . .	5
<b>4</b>	<b>Constraints</b>	<b>6</b>
4.1	GCC Version . . . . .	6
4.2	GDB Version . . . . .	6
<b>5</b>	<b>Guidelines for GCC and Visual C++ compatibility at the same time</b>	<b>7</b>
5.1	Paths . . . . .	7
5.2	Const correctness . . . . .	7
5.3	Templates . . . . .	7
5.4	Member function declarations . . . . .	7
5.5	windows.h . . . . .	8
5.6	Qt is the key . . . . .	8
<b>6</b>	<b>CMake Build System</b>	<b>9</b>
6.1	on Windows . . . . .	9
6.2	on Linux . . . . .	9
6.3	on Mac OS X . . . . .	9
<b>7</b>	<b>Unix fact extraction mechanism</b>	<b>10</b>
7.1	Callmon . . . . .	10
7.2	Patch and Patchclean . . . . .	10
7.3	Metacreator . . . . .	10
<b>8</b>	<b>Tutorial</b>	<b>11</b>
8.1	Preparing the build process . . . . .	11
8.1.1	Pitfalls . . . . .	11
8.2	Building the application . . . . .	11
8.3	Patching the executable . . . . .	11
8.3.1	Patchclean . . . . .	11
8.3.2	Patch . . . . .	11
8.4	Using CGA Toolbar . . . . .	11

8.5	Running the application . . . . .	11
8.6	Using Metacreator . . . . .	11
8.7	Loading the trace(s) into CGA . . . . .	11

# 1 Introduction

## **2 Challenges**

This section describes challenges we had to face while porting the CGA framework to Mac OS X and Linux.

### **2.1 Complexity**

The first challenge was the complexity of CGA. The main application itself contains more than 70.000 lines of code. The whole distribution consists of about 170.000 lines of code.

### **2.2 Platform specifics**

The whole CGA tool chain contains lots of platform specific mechanisms like binary patching, collecting of information from the dynamic linker...

### **2.3 Compiler specifics**

The main instrumentation mechanism is based on a feature provided by the compiler. The Microsoft Visual Studio can insert calls to instrumentation function right after a function was called and right before a function returns. The mechanism in general is the same using gcc, but...

### **2.4 IDE specifics**

While introducing the CMake based build system in CGA...

## **3 Results**

### **3.1 CMake build enviroment**

#### **3.1.1 Qt specific build steps**

### **3.2 Callmon runtime library**

#### **3.2.1 Filesystem events**

#### **3.2.2 Lockfree event to disk**

#### **3.2.3 Asynchronous IO**

### **3.3 Patch and Patchclean**

#### **3.3.1 ELF patching**

#### **3.3.2 MachO patching**

### **3.4 Metacreator**

#### **3.4.1 Line number caching**

## **4 Constraints**

### **4.1 GCC Version**

need  $\geq 4.2$  tested on 4.2, 4.3

### **4.2 GDB Version**

tested on 6.3, 6.8

## 5 Guidelines for GCC and Visual C++ compatibility at the same time

### 5.1 Paths

To be valid on Windows and Unix platforms, a path **must not** contain backslashes as separators. The only valid path separator for both platforms is the slash symbol /. So a valid path looks like this:

```
"this/is/a/valid/path"
```

### 5.2 Const correctness

### 5.3 Templates

The gcc's parser for template type name behaves slightly different than the Visual C++ ones. For example this is a valid definition in Visual C++:

```
std::list<std::pair<int, int>> myListOfIntPairs;
```

This is **not** valid while compiling with gcc. You have to separate > symbols using a space, else, gcc will throw a parser error. So this is the valid equivalent, which compiles on gcc and Visual C++:

```
std::list<std::pair<int, int> > myListOfIntPairs;
```

### 5.4 Member function declarations

When declaring a member function inside the class statement, some people tend to prepend the name of the class to the method name. This may increase readability when inheriting several levels:

```
class A {
public:
    virtual void A::funcFromA();
};

class B : public A {
public:
    virtual void A::funcFromA();
    virtual void B::funcFromB();
};

class C : public B {
public:
    virtual void A::funcFromA();
    virtual void B::funcFromB();
    virtual void C::funcFromC();
};
```

The problem is, this **is not** a valid syntax for gcc. You **must not** prepend the class name to the member function. So the above declaration is valid for gcc like this:

```
class A {
public:
    virtual void funcFromA();
};
```



```
class B : public A {
public:
    virtual void funcFromA();
    virtual void funcFromB();
};

class C : public B {
public:
    virtual void funcFromA();
    virtual void funcFromB();
    virtual void funcFromC();
};
```

## 5.5 windows.h

You **must not** include windows.h because all the types and functions provided by windows.h are highly Windows specific and will not compile nor run on other platforms. In general you will find the same functionality in QtCore. When using QtCore's functionality, it is easy to compile and run the code on other platforms supported by Qt.

## 5.6 Qt is the key

## **6 CMake Build System**

### **6.1 on Windows**

### **6.2 on Linux**

### **6.3 on Mac OS X**

## 7 Unix fact extraction mechanism

### 7.1 Callmon

### 7.2 Patch and Patchclean

### 7.3 Metacreator

## 8 Tutorial

### 8.1 Preparing the build process

#### 8.1.1 Pitfalls

##### Compile with `-finstrument-functions`

**Compile with `-fno-inline`** This disables the inlining of functions which is done by the compiler automatically to optimize execution speed by eliminating function call overhead. Since we are profiling on a function execution level, we cannot profile inlined function, so all the functions inlined by the compiler cannot appear in the call graph. To be sure this cannot happen, use `-fno-inline` as a gcc option. You might skip this if you want. You still might get good profiling results for the calls you are interested in, but you have been warned! Take care!

**Link `unixcallmonlib` as the last library** ...

**On Linux, link with `SYMBOLIC PARAMETER NAME HERE`**

### 8.2 Building the application

### 8.3 Patching the executable

#### 8.3.1 Patchclean

#### 8.3.2 Patch

### 8.4 Using CGA Toolbar

### 8.5 Running the application

### 8.6 Using Metacreator

### 8.7 Loading the trace(s) into CGA